

## Results and Observations (lecture3 lab for COMS 5250)

Results:

==== Performance benchmarks ===

Base N0=20000, repeats=3, steps=4, pop\_fraction=0.5, copy\_batch=50

Average time (seconds)

N	pop_many	pop0_many	L[:]	list(L)	reverse()	[::-1]	x2(pop)	x2(pop0)		
20000	0.000375	0.017067	0.000028	0.000028	0.000034	0.000028	-	-		
40000	0.000693	0.067681	0.000057	0.000054	0.000067	0.000055	1.85x			
80000	0.001403	0.295779	0.000112	0.000109	0.000139	0.000115	2.02x			
160000	0.002972	1.341805	0.000236	0.000236	0.000284	0.000235	2.12x			
							4.37x			
							4.54x			

Observations:

We ran benchmarks for N = 20000, 40000, 80000, 160000. For each N we first generated a base list of length N.

pop\_many means: copy the base list into A, then remove k = N/2 elements using A.pop() (removing from the end) inside a loop, and time the whole operation.

pop0\_many means: copy the base list into A, then remove k = N/2 elements using A.pop(0) (removing from the front) inside a loop, and time the whole operation. L[:] means: time making a copy of the list using slicing (base[:]); list(L) means: time making a copy using the constructor (list(base)); reverse() means: copy the list into A and time A.reverse() which reverses in place; [::-1] means: time base[::-1] which creates a new reversed list...

The measured times for pop\_many were 0.000375 s, 0.000693 s, 0.001403 s, 0.002972 s as N doubled from 20000 to 160000, giving about 1.85x, 2.02x, and 2.12x growth, which likely

is close to linear scaling because popping from the end is  $O(1)$  per operation and we perform  $k = N/2$  removals.

The `pop0_many` times were 0.017067 s, 0.067681 s, 0.295779 s, 1.341805 s, giving about 3.97x, 4.37x, and 4.54x growth when  $N$  doubled. This can probably be considered as consistent with quadratic scaling because each `pop()` forces shifting of the remaining elements, and doing that repeatedly accumulates about  $O(N^2)$  work.

We also observed that as  $N$  increases, `pop0_many` becomes dramatically slower than `pop_many`: at  $N = 20000$  it is  $0.017067/0.000375 \approx 46\times$  slower, and at  $N = 160000$  it is  $1.341805/0.002972 \approx 452\times$  slower. The gap widens quickly because `pop()` repeatedly shifts the remaining elements while `pop()` from the end does not.

For copying, `L[:]` took 0.000028 s, 0.000057 s, 0.000112 s, 0.000236 s and `list(L)` took 0.000028 s, 0.000054 s, 0.000109 s, 0.000236 s, which both scale roughly linearly because both copy  $N$  references into a new list.

For reversing, `reverse()` took 0.000034 s, 0.000067 s, 0.000139 s, 0.000284 s and `[::-1]` took 0.000028 s, 0.000055 s, 0.000115 s, 0.000235 s. Seems like both are roughly linear because both touch all  $N$  elements, but `reverse()` modifies in place while `[::-1]` allocates a new reversed list.